LinCQA: Faster Consistent Query Answering with Linear Time Guarantees

Xiating Ouyang [1]

joint work with Zhiwei Fan [1,2] Paris Koutris [1] Jef Wijsen [3]

University of Wisconsin–Madison [1]

Meta [2]

University of Mons [3]

SIGMOD, Seattle WA, June 18–23 2023

Primary key constraint (violated)

- Metadata of `stackoverflow.com` as of 02/2021 from Stack Exchange Data Dump
- 551M rows, ∼400 GB

Table	# of rows	inRatio	bSize	# of Attributes
Users	14M	0%	1	14
Posts	53M	0%	1	20
PostHistory	141M	0.001%	4	9
Badges	40M	0.58%	941	4
Votes	213M	30.9%	1441	6

inconsistencyRatio = # facts violating PK constraint / # of rows

blockSize = max. # facts with the same PK

Finding consistent answers

Course

c_id	f_id
CS 703	2
CS 703	5
CS 787	3
CS 787	5

CS_Faculty

f_id	f_name
2	Adam
2	Alice
5	Bob

```
SELECT DISTINCT c_id
FROM Course, CS_Faculty
WHERE Course.f_id
      = CS_Faculty.f_id
```

Finding consistent answers

Course

c_id	f_id
CS 703	2
CS 703	5
CS 787	3
CS 787	5

CS_Faculty

f_id	f_name
2	Adam
2	Alice
5	Bob

```
SELECT DISTINCT c_id
FROM Course, CS_Faculty
WHERE Course.f_id
      = CS_Faculty.f_id
```

Return all classes taught by a CS faculty

{CS 703, CS 787} ...

Finding consistent answers

Course

c_id	f_id
CS 703	2
CS 703	5
CS 787	3
CS 787	5

CS_Faculty

f_id	f_name
2	Adam
2	Alice
5	Bob

```
SELECT DISTINCT c_id
FROM Course, CS_Faculty
WHERE Course.f_id
      = CS_Faculty.f_id
```

Return all classes taught by a CS faculty

{CS 703, CS 787} ...

Data cleaning

Finding consistent answers

Course

c_id	f_id
CS 703	2
CS 703	5
CS 787	3
CS 787	5

CS_Faculty

f_id	f_name
2	Adam
2	Alice
5	Bob

```
SELECT DISTINCT c_id
FROM Course, CS_Faculty
WHERE Course.f_id
      = CS_Faculty.f_id
```

Return all classes taught by a CS faculty
{CS 703, CS 787} ...

Data cleaning $2 \times 2 \times 2 \times 1$ repairs

Finding consistent answers

Course

c_id	f_id
CS 703	2
CS 703	5
CS 787	3
CS 787	5

CS_Faculty

f_id	f_name
2	Adam
2	Alice
5	Bob

```
SELECT DISTINCT c_id
FROM Course, CS_Faculty
WHERE Course.f_id
      = CS_Faculty.f_id
```

Return all classes taught by a CS faculty
{CS 703, CS 787} ...

Data cleaning $2 \times 2 \times 2 \times 1$ repairs

What are the answers guaranteed to be returned on all repairs?

Finding consistent answers

Course

c_id	f_id
CS 703	2
CS 703	5
CS 787	3
CS 787	5

CS_Faculty

f_id	f_name
2	Adam
2	Alice
5	Bob

```
SELECT DISTINCT c_id
FROM Course, CS_Faculty
WHERE Course.f_id
      = CS_Faculty.f_id
```

Return all classes taught by a CS faculty
{CS 703, CS 787} ...

Data cleaning $2 \times 2 \times 2 \times 1$ repairs

What are the answers guaranteed to be returned on all repairs?

CS 703

Finding consistent answers

Course

c_id	f_id
CS 703	2
CS 703	5
CS 787	3
CS 787	5

CS_Faculty

f_id	f_name
2	Adam
2	Alice
5	Bob

```
SELECT DISTINCT c_id
FROM Course, CS_Faculty
WHERE Course.f_id
      = CS_Faculty.f_id
```

Return all classes taught by a CS faculty
{CS 703, CS 787} ...

Data cleaning $2 \times 2 \times 2 \times 1$ repairs

What are the answers guaranteed to be returned on all repairs?

CS 703

Consistent Answer

Finding consistent answers without enumeration

Course

c_id	f_id
CS 703	2
CS 703	5
CS 787	3
CS 787	5

CS_Faculty

f_id	f_name
2	Adam
2	Alice
5	Bob

```
SELECT DISTINCT c_id
FROM Course, CS_Faculty
WHERE Course.f_id
      = CS_Faculty.f_id  AND
  (all f_id's for the same c_id
  appear in CS_Faculty)
```

The original query has a first-order rewriting

Finding consistent answers without enumeration

Course

c_id	f_id
CS 703	2
CS 703	5
CS 787	3
CS 787	5

CS_Faculty

f_id	f_name
2	Adam
2	Alice
5	Bob

```
SELECT DISTINCT c_id
FROM Course, CS_Faculty
WHERE Course.f_id
      = CS_Faculty.f_id  AND
  (all f_id's for the same c_id
  appear in CS_Faculty)
```

The original query has a first-order rewriting

Finding consistent answers without enumeration

Course

c_id	f_id
CS 703	2
CS 703	5
CS 787	3
CS 787	5

CS_Faculty

f_id	f_name
2	Adam
2	Alice
5	Bob

```
SELECT DISTINCT c_id
FROM Course, CS_Faculty
WHERE Course.f_id
      = CS_Faculty.f_id  AND
  (all f_id's for the same c_id
  appear in CS_Faculty)
```

The original query has a first-order rewriting

Finding consistent answers without enumeration

Course

c_id	f_id
CS 703	2
CS 703	5
CS 787	3
CS 787	5

CS_Faculty

f_id	f_name
2	Adam
2	Alice
5	Bob

```
SELECT DISTINCT c_id
FROM Course, CS_Faculty
WHERE Course.f_id
      = CS_Faculty.f_id  AND
  (all f_id's for the same c_id
  appear in CS_Faculty)
```

The original query has a first-order rewriting

For which Q can the consistent answers be found efficiently?

Can we build a system to find the consistent answers?

For which Q can the consistent answers be found efficiently?

Can we build a system to find the consistent answers?

Acyclic query evaluation

$q() \text{ :- } \text{Course}(\underline{x}, y), \text{CS\_Faculty}(\underline{y}, z)$

Acyclic query evaluation

$q() \text{ :- } \text{Course}(\underline{x}, y), \text{CS\_Faculty}(\underline{y}, z)$

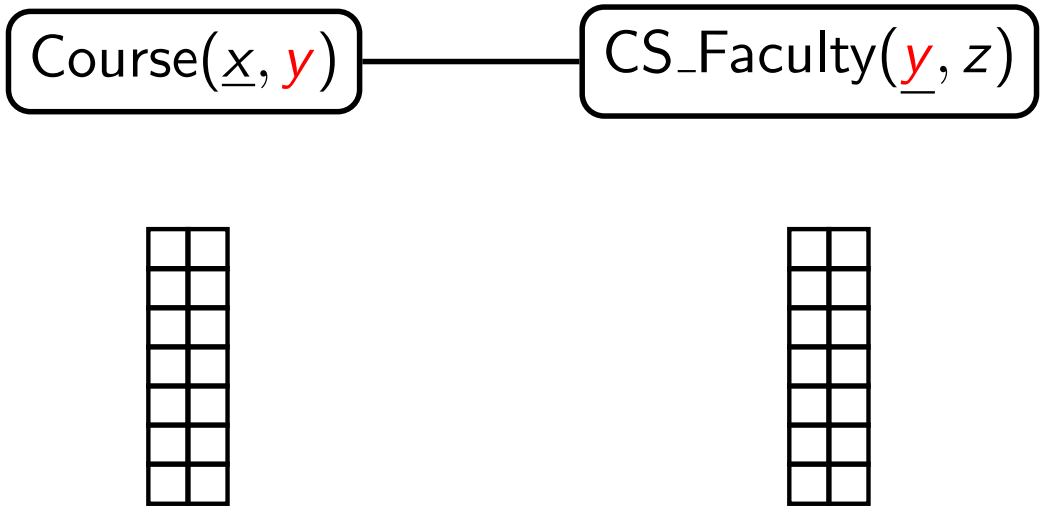

Acyclic query evaluation

$q() \text{ :- } \text{Course}(\underline{x}, y), \text{CS\_Faculty}(\underline{y}, z)$

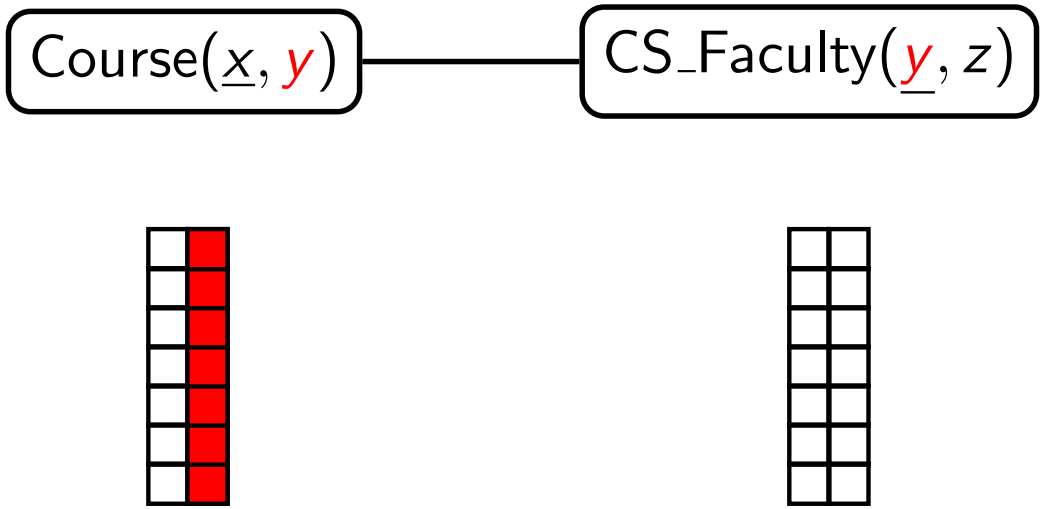

Acyclic query evaluation

$q() \text{ :- } \text{Course}(\underline{x}, y), \text{CS\_Faculty}(\underline{y}, z)$

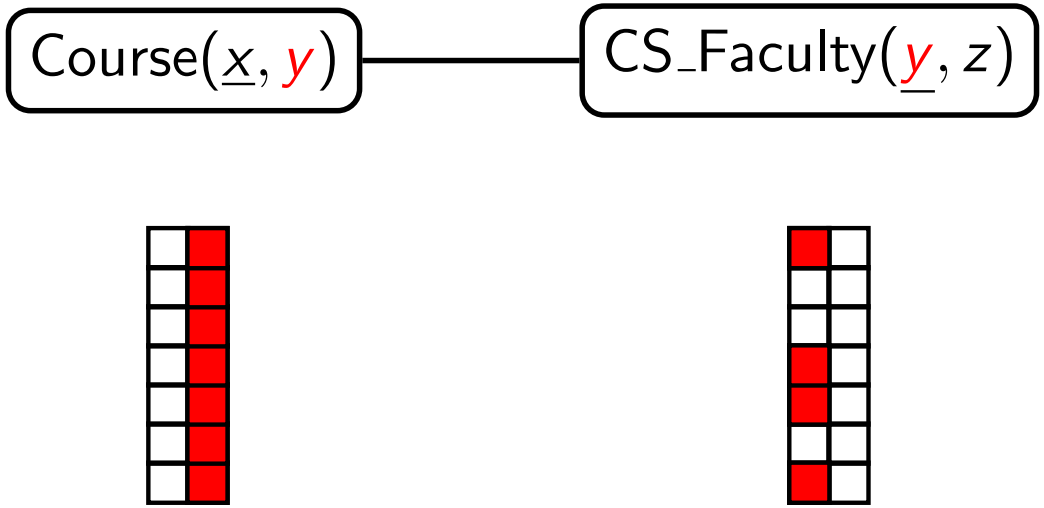

Yannakakis [VLDB'81]

The answer to every **Boolean** acyclic query can be computed in $O(|\mathbf{db}|)$.

Yannakakis [VLDB'81] Our result

consistent answer

The ~~answer~~ to every **Boolean** acyclic query can be computed in $O(|\mathbf{db}|)$.

$\wedge$

with a pair-pruning join tree (PPJT)

Yannakakis [VLDB'81] Our result

consistent answer

The ~~answer~~ to every **Boolean** acyclic query can be computed in $O(|\mathbf{db}|)$.

$\wedge$

with a pair-pruning join tree (PPJT)

non-Boolean $\leq_T^P$ **Boolean**

Pair-pruning join tree (PPJT)

A join tree rooted at some atom is a PPJT if

the root of every subtree is unattacked in the subtree

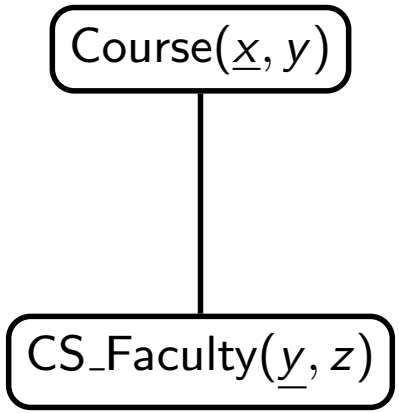

Pair-pruning join tree (PPJT)

A join tree rooted at some atom is a PPJT if

the root of every subtree is unattacked in the subtree

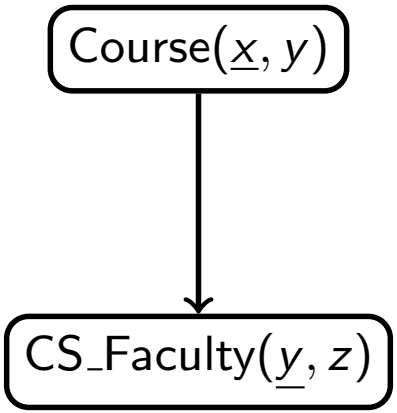

Pair-pruning join tree (PPJT)

A join tree rooted at some atom is a PPJT if

the root of every subtree is unattacked in the subtree

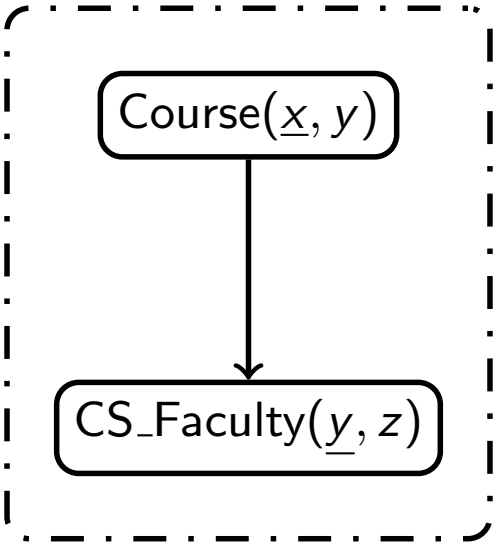

Pair-pruning join tree (PPJT)

A join tree rooted at some atom is a PPJT if

the root of every subtree is unattacked in the subtree

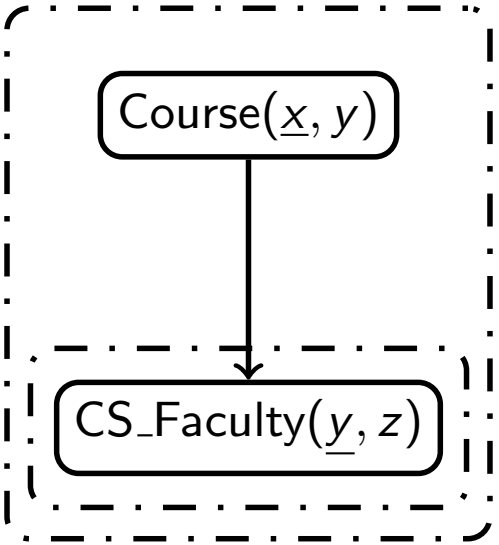

Pair-pruning join tree (PPJT)

A join tree rooted at some atom is a PPJT if

the root of every subtree is unattacked in the subtree

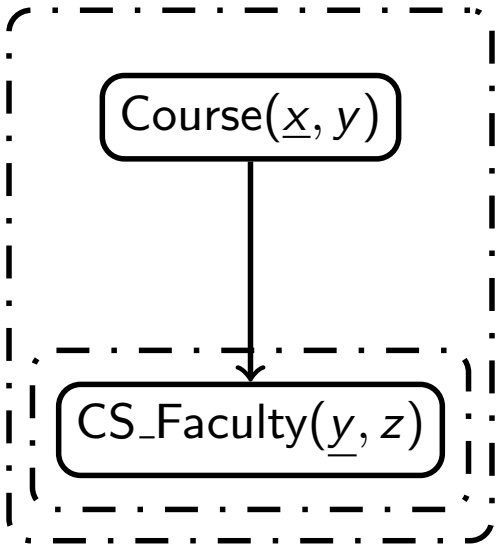

Every acyclic query has a join tree, but not every acyclic query has a PPJT

PPJT is a wide class

+ star/snowflake schema (e.g. TPC-H)
+ Every acyclic query in $\mathcal{C}_{\text{forest}}$ [FM, ICDT'05] has a PPJT

From PPJT to **FO**-rewriting

Remove a primary key if some tuple with this primary key is "bad"

$\text{Course}_{join}() \text{ :- } \text{Course}(x, y), \neg\text{Course}_{fkey}(x)$

$\text{Course}_{fkey}(x) \text{ :- } \text{Course}(x, y), \neg\text{Faculty}_{join}(y)$

$\forall\text{Child} : \text{Root}_{fkey}(\vec{x}) \text{ :- } \text{Root}(\vec{x}, \vec{y}), \neg\text{Child}_{join}(\vec{\alpha})$

$\text{Child}_{join}(\vec{\alpha}) \text{ :- } \text{Child}(\vec{u}, \vec{v}), \neg\text{Child}_{fkey}(\vec{u})$

$\text{Faculty}_{join}(y) \text{ :- } \text{Faculty}(y, z, w), \neg\text{Faculty}_{fkey}(y)$

$\text{Faculty}_{fkey}(y) \text{ :- } \text{Faculty}(y, z, w), w \neq \text{"DB"}$

also expressible in SQL!
runs in $O(N)$

From PPJT to **FO**-rewriting

Remove a primary key if some tuple with this primary key is "bad"

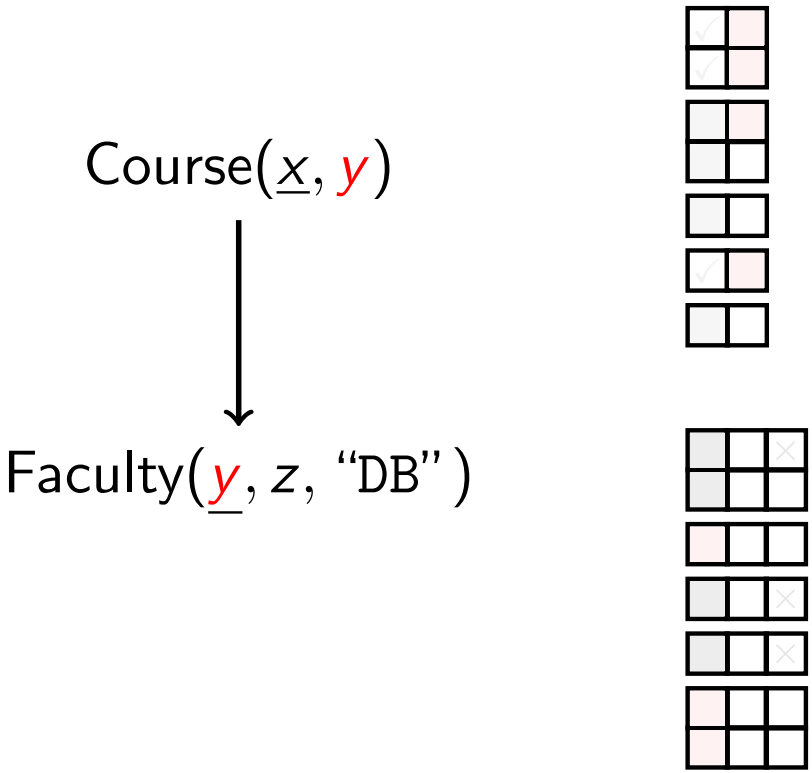

$\text{Course}_{join}() \text{ :- } \text{Course}(x, y), \neg\text{Course}_{fkey}(x)$

$\text{Course}_{fkey}(x) \text{ :- } \text{Course}(x, y), \neg\text{Faculty}_{join}(y)$

$\forall\text{Child} : \text{Root}_{fkey}(\vec{x}) \text{ :- } \text{Root}(\vec{x}, \vec{y}), \neg\text{Child}_{join}(\vec{\alpha})$

$\text{Child}_{join}(\vec{\alpha}) \text{ :- } \text{Child}(\vec{u}, \vec{v}), \neg\text{Child}_{fkey}(\vec{u})$

$\text{Faculty}_{join}(y) \text{ :- } \text{Faculty}(y, z, w), \neg\text{Faculty}_{fkey}(y)$

$\text{Faculty}_{fkey}(y) \text{ :- } \text{Faculty}(y, z, w), w \neq \text{"DB"}$

also expressible in SQL!
runs in $O(N)$

From PPJT to **FO**-rewriting

Remove a primary key if some tuple with this primary key is "bad"

Course($\underline{x}$, y)

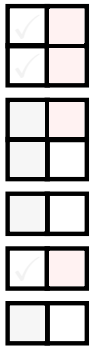

Faculty($\underline{y}$, z, "DB")

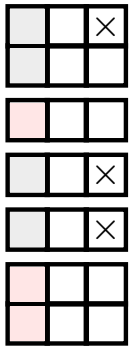

$\text{Course}_{join}() \text{ :- } \text{Course}(x, y), \neg\text{Course}_{fkey}(x)$

$\text{Course}_{fkey}(x) \text{ :- } \text{Course}(x, y), \neg\text{Faculty}_{join}(y)$

$\forall\text{Child} : \text{Root}_{fkey}(\vec{x}) \text{ :- } \text{Root}(\underline{\vec{x}}, \vec{y}), \neg\text{Child}_{join}(\vec{\alpha})$

$\text{Child}_{join}(\vec{\alpha}) \text{ :- } \text{Child}(\underline{\vec{u}}, \vec{v}), \neg\text{Child}_{fkey}(\vec{u})$

$\text{Faculty}_{join}(y) \text{ :- } \text{Faculty}(y, z, w), \neg\text{Faculty}_{fkey}(y)$

$\text{Faculty}_{fkey}(y) \text{ :- } \text{Faculty}(y, z, w), w \neq \text{"DB"}$

also expressible in SQL!
runs in $O(N)$

From PPJT to **FO**-rewriting

Remove a primary key if some tuple with this primary key is "bad"

Course($\underline{x}$, y)

Faculty($\underline{y}$, z, "DB")

$\text{Course}_{join}() \text{ :- } \text{Course}(x, y), \neg\text{Course}_{fkey}(x)$

$\text{Course}_{fkey}(x) \text{ :- } \text{Course}(x, y), \neg\text{Faculty}_{join}(y)$

$\forall\text{Child} : \text{Root}_{fkey}(\vec{x}) \text{ :- } \text{Root}(\underline{\vec{x}}, \vec{y}), \neg\text{Child}_{join}(\vec{\alpha})$

$\text{Child}_{join}(\vec{\alpha}) \text{ :- } \text{Child}(\underline{\vec{u}}, \vec{v}), \neg\text{Child}_{fkey}(\vec{u})$

$\text{Faculty}_{join}(y) \text{ :- } \text{Faculty}(y, z, w), \neg\text{Faculty}_{fkey}(y)$

$\text{Faculty}_{fkey}(y) \text{ :- } \text{Faculty}(y, z, w), w \neq \text{"DB"}$

also expressible in SQL!
runs in $O(N)$

From PPJT to **FO**-rewriting

Remove a primary key if some tuple with this primary key is "bad"

Course($\underline{x}$, y)

↓

Faculty($\underline{y}$, z, "DB")

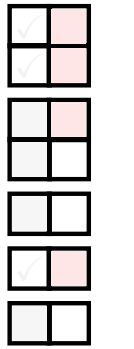

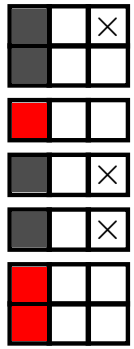

$\text{Course}_{join}() \text{ :- } \text{Course}(x, y), \neg\text{Course}_{fkey}(x)$

$\text{Course}_{fkey}(x) \text{ :- } \text{Course}(x, y), \neg\text{Faculty}_{join}(y)$

$\forall\text{Child} : \text{Root}_{fkey}(\vec{x}) \text{ :- } \text{Root}(\underline{\vec{x}}, \vec{y}), \neg\text{Child}_{join}(\vec{\alpha})$

$\text{Child}_{join}(\vec{\alpha}) \text{ :- } \text{Child}(\underline{\vec{u}}, \vec{v}), \neg\text{Child}_{fkey}(\vec{u})$

$\text{Faculty}_{join}(y) \text{ :- } \text{Faculty}(y, z, w), \neg\text{Faculty}_{fkey}(y)$

$\text{Faculty}_{fkey}(y) \text{ :- } \text{Faculty}(y, z, w), w \neq \text{"DB"}$

also expressible in SQL!

runs in $O(N)$

From PPJT to **FO**-rewriting

Remove a primary key if some tuple with this primary key is "bad"

$\text{Course}(\underline{x}, y)$

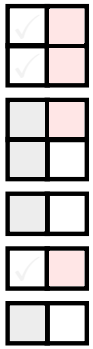

$\text{Faculty}(\underline{y}, z, \text{"DB"})$

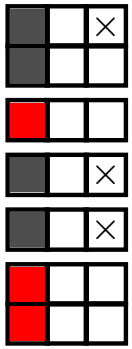

$\text{Course}_{join}() \text{ :- } \text{Course}(x, y), \neg\text{Course}_{fkey}(x)$

$\text{Course}_{fkey}(x) \text{ :- } \text{Course}(x, y), \neg\text{Faculty}_{join}(y)$

$\forall \text{Child} : \text{Root}_{fkey}(\vec{x}) \text{ :- } \text{Root}(\underline{\vec{x}}, \vec{y}), \neg\text{Child}_{join}(\vec{\alpha})$

$\boxed{\text{Child}_{join}(\vec{\alpha}) \text{ :- } \text{Child}(\underline{\vec{u}}, \vec{v}), \neg\text{Child}_{fkey}(\vec{u})}$

$\text{Faculty}_{join}(y) \text{ :- } \text{Faculty}(y, z, w), \neg\text{Faculty}_{fkey}(y)$

$\text{Faculty}_{fkey}(y) \text{ :- } \text{Faculty}(y, z, w), w \neq \text{"DB"}$

also expressible in SQL!
runs in $O(N)$

From PPJT to **FO**-rewriting

Remove a primary key if some tuple with this primary key is "bad"

$\mathsf{Course}(\underline{x}, y)$

$\mathsf{Faculty}(\underline{y}, z, \text{“DB”})$

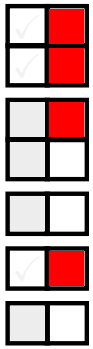

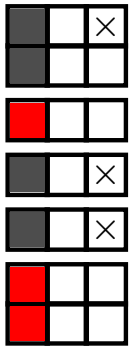

$\mathsf{Course}_{join}() \text{ :- } \mathsf{Course}(x, y), \neg\mathsf{Course}_{fkey}(x)$

$\mathsf{Course}_{fkey}(x) \text{ :- } \mathsf{Course}(x, y), \neg\mathsf{Faculty}_{join}(y)$

$\forall\mathsf{Child} : \mathsf{Root}_{fkey}(\vec{x}) \text{ :- } \mathsf{Root}(\underline{\vec{x}}, \vec{y}), \neg\mathsf{Child}_{join}(\vec{\alpha})$

$\boxed{\mathsf{Child}_{join}(\vec{\alpha}) \text{ :- } \mathsf{Child}(\underline{\vec{u}}, \vec{v}), \neg\mathsf{Child}_{fkey}(\vec{u})}$

$\mathsf{Faculty}_{join}(y) \text{ :- } \mathsf{Faculty}(y, z, w), \neg\mathsf{Faculty}_{fkey}(y)$

$\mathsf{Faculty}_{fkey}(y) \text{ :- } \mathsf{Faculty}(y, z, w), w \neq \text{“DB”}$

also expressible in SQL!
runs in $O(N)$

From PPJT to **FO**-rewriting

Remove a primary key if some tuple with this primary key is "bad"

Course($\underline{x}$, y)

↓

Faculty($\underline{y}$, z, "DB")

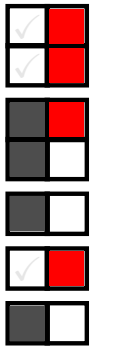

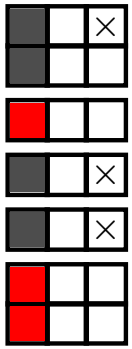

$\text{Course}_{join}() \text{ :- } \text{Course}(x, y), \neg\text{Course}_{fkey}(x)$

$\text{Course}_{fkey}(x) \text{ :- } \text{Course}(x, y), \neg\text{Faculty}_{join}(y)$

$\forall\text{Child} : \text{Root}_{fkey}(\vec{x}) \text{ :- } \text{Root}(\underline{\vec{x}}, \vec{y}), \neg\text{Child}_{join}(\vec{\alpha})$

$\boxed{\text{Child}_{join}(\vec{\alpha}) \text{ :- } \text{Child}(\underline{\vec{u}}, \vec{v}), \neg\text{Child}_{fkey}(\vec{u})}$

$\text{Faculty}_{join}(y) \text{ :- } \text{Faculty}(y, z, w), \neg\text{Faculty}_{fkey}(y)$

$\text{Faculty}_{fkey}(y) \text{ :- } \text{Faculty}(y, z, w), w \neq \text{"DB"}$

also expressible in SQL!
runs in $O(N)$

From PPJT to **FO**-rewriting

Remove a primary key if some tuple with this primary key is "bad"

Course($\underline{x}$, y)

↓

Faculty($\underline{y}$, z, "DB")

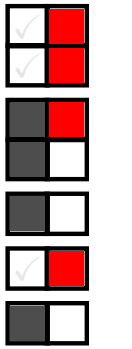

$\text{Course}_{join}() \text{ :- } \text{Course}(x, y), \neg\text{Course}_{fkey}(x)$

$\text{Course}_{fkey}(x) \text{ :- } \text{Course}(x, y), \neg\text{Faculty}_{join}(y)$

$\boxed{\forall\text{Child} : \text{Root}_{fkey}(\vec{x}) \text{ :- } \text{Root}(\underline{\vec{x}}, \vec{y}), \neg\text{Child}_{join}(\vec{\alpha})}$

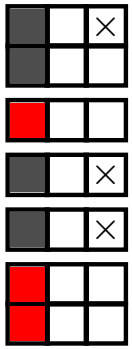

$\boxed{\text{Child}_{join}(\vec{\alpha}) \text{ :- } \text{Child}(\underline{\vec{u}}, \vec{v}), \neg\text{Child}_{fkey}(\vec{u})}$

$\text{Faculty}_{join}(y) \text{ :- } \text{Faculty}(y, z, w), \neg\text{Faculty}_{fkey}(y)$

$\text{Faculty}_{fkey}(y) \text{ :- } \text{Faculty}(y, z, w), w \neq \text{"DB"}$

also expressible in SQL!
runs in $O(N)$

From PPJT to **FO**-rewriting

Remove a primary key if some tuple with this primary key is "bad"

Course($\underline{x}$, y)

↓

Faculty($\underline{y}$, z, "DB")

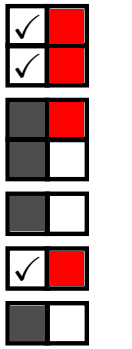

$\text{Course}_{join}() \text{ :- } \text{Course}(x, y), \neg\text{Course}_{fkey}(x)$

$\text{Course}_{fkey}(x) \text{ :- } \text{Course}(x, y), \neg\text{Faculty}_{join}(y)$

$\boxed{\forall\text{Child} : \text{Root}_{fkey}(\vec{x}) \text{ :- } \text{Root}(\underline{\vec{x}}, \vec{y}), \neg\text{Child}_{join}(\vec{\alpha})}$

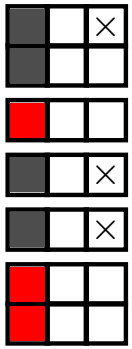

$\boxed{\text{Child}_{join}(\vec{\alpha}) \text{ :- } \text{Child}(\underline{\vec{u}}, \vec{v}), \neg\text{Child}_{fkey}(\vec{u})}$

$\text{Faculty}_{join}(y) \text{ :- } \text{Faculty}(y, z, w), \neg\text{Faculty}_{fkey}(y)$

$\text{Faculty}_{fkey}(y) \text{ :- } \text{Faculty}(y, z, w), w \neq \text{"DB"}$

also expressible in SQL!
runs in $O(N)$

From PPJT to **FO**-rewriting

Remove a primary key if some tuple with this primary key is "bad"

Course($\underline{x}$, y)

↓

Faculty($\underline{y}$, z, "DB")

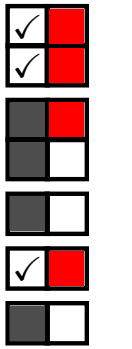

$\text{Course}_{join}() \text{ :- } \text{Course}(x, y), \neg\text{Course}_{fkey}(x)$

$\text{Course}_{fkey}(x) \text{ :- } \text{Course}(x, y), \neg\text{Faculty}_{join}(y)$

$\boxed{\forall\text{Child} : \text{Root}_{fkey}(\vec{x}) \text{ :- } \text{Root}(\underline{\vec{x}}, \vec{y}), \neg\text{Child}_{join}(\vec{\alpha})}$

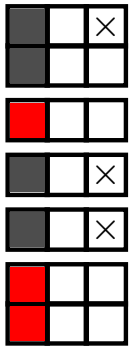

$\boxed{\text{Child}_{join}(\vec{\alpha}) \text{ :- } \text{Child}(\underline{\vec{u}}, \vec{v}), \neg\text{Child}_{fkey}(\vec{u})}$

$\text{Faculty}_{join}(y) \text{ :- } \text{Faculty}(y, z, w), \neg\text{Faculty}_{fkey}(y)$

$\text{Faculty}_{fkey}(y) \text{ :- } \text{Faculty}(y, z, w), w \neq \text{"DB"}$

also expressible in SQL!

runs in $O(N)$

From PPJT to **FO**-rewriting

Remove a primary key if some tuple with this primary key is "bad"

Course($\underline{x}$, y)

↓

Faculty($\underline{y}$, z, "DB")

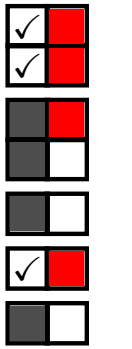

$\text{Course}_{join}() \text{ :- } \text{Course}(x, y), \neg\text{Course}_{fkey}(x)$

$\text{Course}_{fkey}(x) \text{ :- } \text{Course}(x, y), \neg\text{Faculty}_{join}(y)$

$\boxed{\forall\text{Child} : \text{Root}_{fkey}(\vec{x}) \text{ :- } \text{Root}(\underline{\vec{x}}, \vec{y}), \neg\text{Child}_{join}(\vec{\alpha})}$

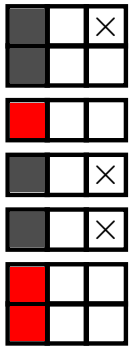

$\boxed{\text{Child}_{join}(\vec{\alpha}) \text{ :- } \text{Child}(\underline{\vec{u}}, \vec{v}), \neg\text{Child}_{fkey}(\vec{u})}$

$\text{Faculty}_{join}(y) \text{ :- } \text{Faculty}(y, z, w), \neg\text{Faculty}_{fkey}(y)$

$\text{Faculty}_{fkey}(y) \text{ :- } \text{Faculty}(y, z, w), w \neq \text{"DB"}$

also expressible in SQL!

runs in $O(N)$

From Boolean to non-Boolean

```
SELECT DISTINCT A1, A2 FROM T WHERE A3 = 42
```

Step 1 Evaluate directly

A1	A2
a	b
x	y
...	...

Step 2 Reduce to **Boolean** (using PPJT)

```
SELECT DISTINCT 1 FROM T WHERE A3 = 42 AND A1 = a AND A2 = b
```

if **yes**, then output (a, b), otherwise continue

```
SELECT DISTINCT 1 FROM T WHERE A3 = 42 AND A1 = x AND A2 = y
```

...

$\xrightarrow{\text{LinCQA}}$ a single SQL/Datalog query

From Boolean to non-Boolean

```
SELECT DISTINCT A1, A2 FROM T WHERE A3 = 42
```

Step 1 Evaluate directly

A1	A2
a	b
x	y
...	...

Step 2 Reduce to **Boolean** (using PPJT)

```
SELECT DISTINCT 1 FROM T WHERE A3 = 42 AND A1 = a AND A2 = b
```

if **yes**, then output (a, b), otherwise continue

```
SELECT DISTINCT 1 FROM T WHERE A3 = 42 AND A1 = x AND A2 = y
```

...

$\xrightarrow{\text{LinCQA}}$ a single SQL/Datalog query

From Boolean to non-Boolean

```
SELECT DISTINCT A1, A2 FROM T WHERE A3 = 42
```

Step 1 Evaluate directly

A1	A2
a	b
x	y
...	...

Step 2 Reduce to **Boolean** (using PPJT)

```
SELECT DISTINCT 1 FROM T WHERE A3 = 42 AND A1 = a AND A2 = b
```

if **yes**, then output (a, b), otherwise continue

```
SELECT DISTINCT 1 FROM T WHERE A3 = 42 AND A1 = x AND A2 = y
```

...

$\xrightarrow{\text{LinCQA}}$ a single SQL/Datalog query

From Boolean to non-Boolean

```
SELECT DISTINCT A1, A2 FROM T WHERE A3 = 42
```

Step 1 Evaluate directly

A1	A2
a	b
x	y
...	...

Step 2 Reduce to **Boolean** (using PPJT)

```
SELECT DISTINCT 1 FROM T WHERE A3 = 42 AND A1 = a AND A2 = b
```

if **yes**, then output (a, b), otherwise continue

```
SELECT DISTINCT 1 FROM T WHERE A3 = 42 AND A1 = x AND A2 = y
```

...

$\xrightarrow{\text{LinCQA}}$ a single SQL/Datalog query

From Boolean to non-Boolean

```
SELECT DISTINCT A1, A2 FROM T WHERE A3 = 42
```

Step 1 Evaluate directly

A1	A2
a	b
x	y
...	...

Step 2 Reduce to **Boolean** (using PPJT)

```
SELECT DISTINCT 1 FROM T WHERE A3 = 42 AND A1 = a AND A2 = b
```

if **yes**, then output (a, b), otherwise continue

```
SELECT DISTINCT 1 FROM T WHERE A3 = 42 AND A1 = x AND A2 = y
```

...

$\xrightarrow{\text{LinCQA}}$ a single SQL/Datalog query

Acyclic q	PPJT	Yannakakis [VLDB'81]
Boolean q	$O(N)$	$O(N)$
non-Boolean q	$O(N \cdot \lvert\mathsf{OUT}_{\mathsf{inconsistent}}\rvert)$	$O(N \cdot \lvert\mathsf{OUT}\rvert)$
full q (`SELECT *`)	$O(N + \lvert\mathsf{OUT}_{\mathsf{consistent}}\rvert)$	$O(N + \lvert\mathsf{OUT}\rvert)$

Consistent answers of common join queries can be computed with no asymptotic overhead

Acyclic q	PPJT	Yannakakis [VLDB'81]
Boolean q	$O(N)$	$O(N)$
non-Boolean q	$O(N \cdot \lvert\text{OUT}_{\text{inconsistent}}\rvert)$	$O(N \cdot \lvert\text{OUT}\rvert)$
full q (`SELECT *`)	$O(N + \lvert\text{OUT}_{\text{consistent}}\rvert)$	$O(N + \lvert\text{OUT}\rvert)$

Consistent answers of common join queries can be computed with no asymptotic overhead

Setup & Baselines

System	Target class	Interm. output	Backend
CAvSAT	*	SAT formula	SQL Server & MaxHS
Conquer	$\mathcal{C}_{\text{forest}}$	SQL	SQL Server
Improved Conquesto	SJF **FO**	SQL	SQL Server
LinCQA	PPJT	SQL	SQL Server

Stackoverflow data

- Metadata of `stackoverflow.com` as of 02/2021 from Stack Exchange Data Dump
- 551M rows, 400 GB

Table	# of rows	inRatio	bSize	# of Attributes
Users	14M	0%	1	14
Posts	53M	0%	1	20
PostHistory	141M	0.001%	4	9
Badges	40M	0.58%	941	4
Votes	213M	30.9%	1441	6

Experiments on Stackoverflow

Q_1 : Posts $\bowtie$ Votes Q_2 : Users $\bowtie$ Badges Q_3 : Users $\bowtie$ Posts

Q_4 : Users $\bowtie$ Posts $\bowtie$ Comments

Q_5 : Posts $\bowtie$ PostHistory $\bowtie$ Votes $\bowtie$ Comments

```
SELECT DISTINCT Posts.Title
FROM Posts, PostHistory, Votes, Comments
WHERE Posts.Tags LIKE "%SQL%"
    AND Posts.id = PostHistory.PostId
    AND Posts.id = Comments.PostId
    AND Posts.id = Votes.PostId
    AND Votes.BountyAmount > 100
    AND PostHistory.PostHistoryTypeId = 2
    AND Comments.score = 0
```

```
    )
),

PostHistory_bad_key as (
    select PH.PostId, PH.CreationDate, PH.UserId, PH.PostHistoryTypeId
    from PostHistory PH
    where PH.PostHistoryTypeId <> 2
),

PostHistory_good_join as (
    select PH.PostId
    from PostHistory PH
    where not exists (
        select *
        from PostHistory_bad_key
        where PH.PostId = PostHistory_bad_key.PostId and
            PH.CreationDate = PostHistory_bad_key.CreationDate and
            PH.UserId = PostHistory_bad_key.UserId and
            PH.PostHistoryTypeId = PostHistory_bad_key.PostHistoryTypeId
    )
),

Votes_bad_key as (
    select V.PostId, V.UserId, V.CreationDate
    from Votes V
    where V.BountyAmount <= 100 or V.BountyAmount is null
),

Votes_good_join as (
    select V.PostId
    from Votes V
    where not exists (
        select *
        from Votes_bad_key
        where
        V.PostId = Votes_bad_key.PostId and
        V.UserId = Votes_bad_key.UserId and
        V.CreationDate = Votes_bad_key.CreationDate
    )
),

Comments_bad_key as (
    select C.CreationDate, C.UserId, candidates.Title
    from Comments C
    join candidates on (C.CreationDate = candidates.CreationDate and C.UserId = candidates.UserId)
    where C.score <> 0


    union all

    select C.CreationDate, C.UserId, candidates.Title
    from Comments C
    join candidates on (C.CreationDate = candidates.CreationDate and C.UserId = candidates.UserId)
    left outer join Posts_good_join on (C.PostId = Posts_good_join.Id and candidates.Title = Posts_good_join.Title)
    left outer join PostHistory_good_join on (C.PostId = PostHistory_good_join.PostId)
    left outer join Votes_good_join on (C.PostId = Votes_good_join.PostId)
    where (
        Posts_good_join.Id is NULL or PostHistory_good_join.PostId is NULL or Votes_good_join.PostId is NULL
        or Posts_good_join.Title is NULL
    )
),

Comments_good_join as (
    select candidates.Title
    from Comments C
    join candidates on (C.CreationDate = candidates.CreationDate and C.UserId = candidates.UserId)
    where not exists (
        select *
        from Comments_bad_key
        where C.CreationDate = Comments_bad_key.CreationDate
        and C.UserId = Comments_bad_key.UserId
        and candidates.Title = Comments_bad_key.Title
    )
)

select distinct Title
from Comments_good_join
```

Experiments on Stackoverflow

Q_1 : Posts $\bowtie$ Votes Q_2 : Users $\bowtie$ Badges Q_3 : Users $\bowtie$ Posts

Q_4 : Users $\bowtie$ Posts $\bowtie$ Comments

Q_5 : Posts $\bowtie$ PostHistory $\bowtie$ Votes $\bowtie$ Comments

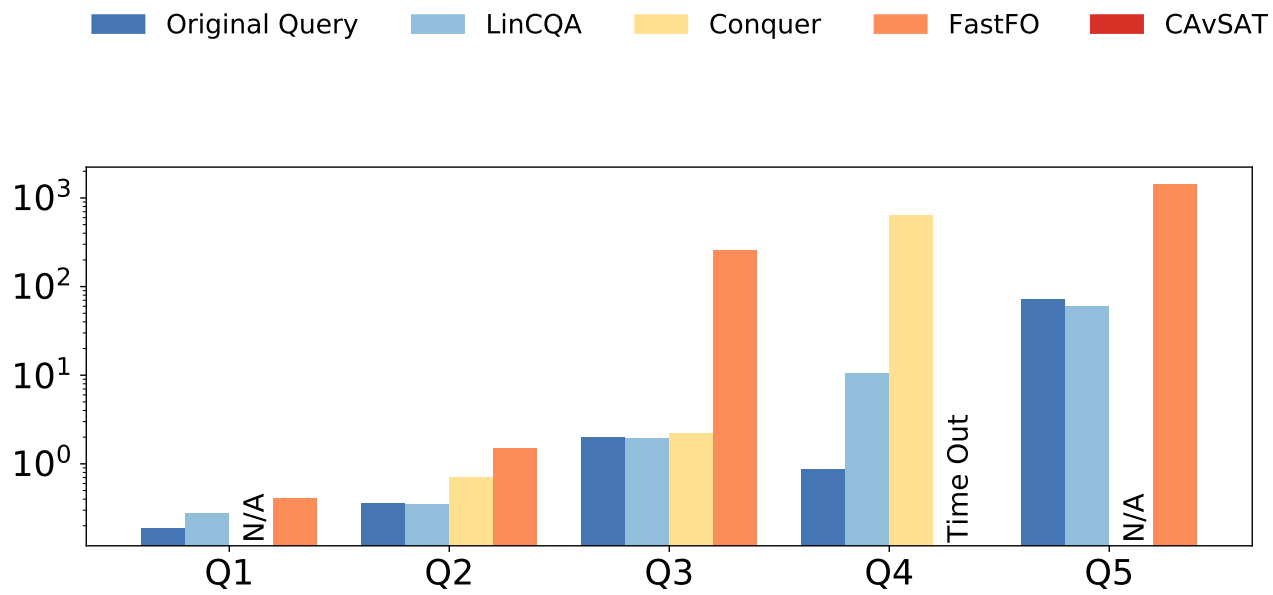

# poss.	27578	145	38320	3925	1250
# cons.	27578	145	38320	3925	1245

Concluding remarks

Acyclic q	PPJT	Yannakakis [VLDB'81]
Boolean q	$O(N)$	$O(N)$
non-Boolean q	$O(N \cdot \lvert \mathsf{OUT}_{\mathsf{inconsistent}} \rvert)$	$O(N \cdot \lvert \mathsf{OUT} \rvert)$
full q (`SELECT *`)	$O(N + \lvert \mathsf{OUT}_{\mathsf{consistent}} \rvert)$	$O(N + \lvert \mathsf{OUT} \rvert)$

Concluding remarks

Acyclic q	PPJT	Yannakakis [VLDB'81]
Boolean q	$O(N)$	$O(N)$
non-Boolean q	$O(N \cdot \lvert\text{OUT}_{\text{inconsistent}}\rvert)$	$O(N \cdot \lvert\text{OUT}\rvert)$
full q (`SELECT *`)	$O(N + \lvert\text{OUT}_{\text{consistent}}\rvert)$	$O(N + \lvert\text{OUT}\rvert)$

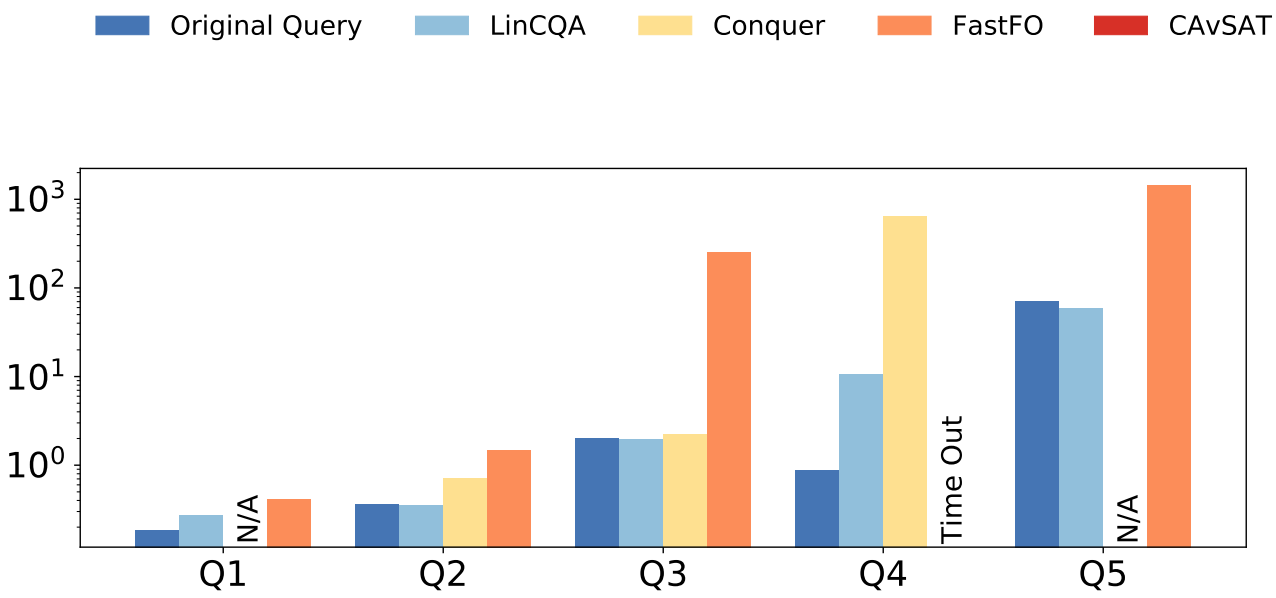

Concluding remarks

Acyclic q	PPJT	Yannakakis [VLDB'81]
Boolean q	$O(N)$	$O(N)$
non-Boolean q	$O(N \cdot \lvert \mathrm{OUT}_{\mathrm{inconsistent}} \rvert)$	$O(N \cdot \lvert \mathrm{OUT} \rvert)$
full q (`SELECT *`)	$O(N + \lvert \mathrm{OUT}_{\mathrm{consistent}} \rvert)$	$O(N + \lvert \mathrm{OUT} \rvert)$

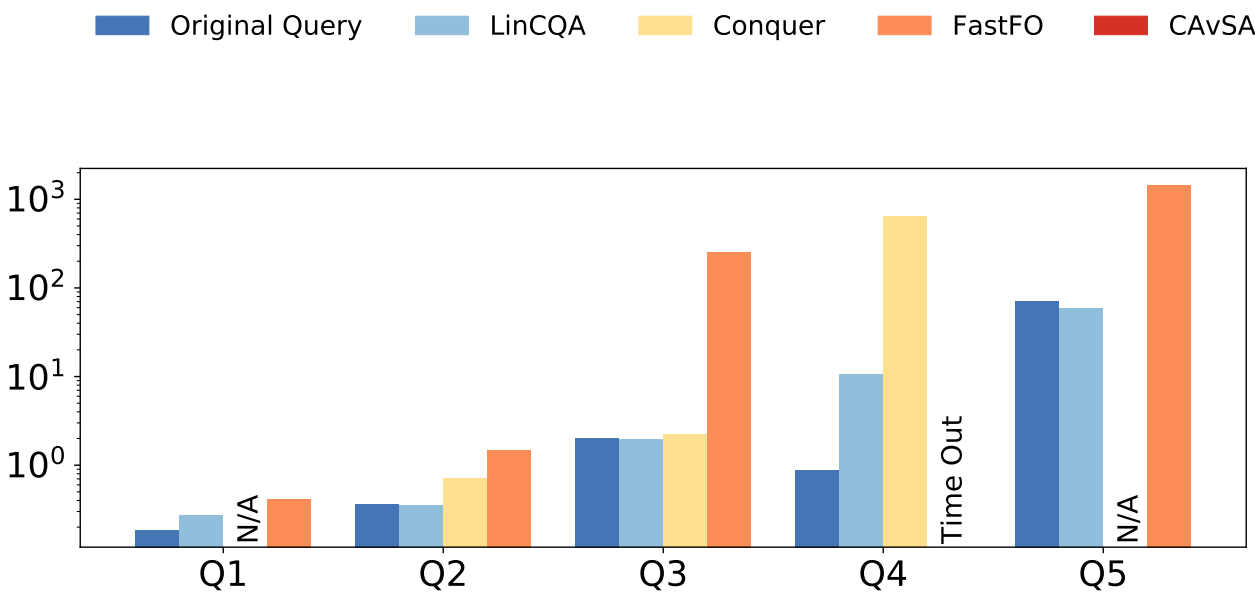

Thank you!